

# Cleaning Data - due thu 09/30 | *by Charlotte*

---

## Exercise 1

For our data cleaning exercises, we will return one last time to our ACS data here. Download and import the 10percent ACS sample.

```
In [ ]: import pandas as pd
        acs = pd.read_stata('US_ACS_2017_10pct_sample.dta')
```

---

## Exercise 2

For our exercises today, we'll focus on age, gender, sex, and inctot. Subset your data to those variables, and quickly look at a sample of 10 rows.

```
In [ ]: sub_acs = acs.loc[:, (acs.columns == 'age') | (acs.columns == 'sex') | (acs
sub_acs
import numpy.random as npr
npr.seed(42)
sub_acs.sample(10)
```

```
Out[ ]:
```

	sex	age	educ	inctot
166590	male	62	1 year of college	170000
207895	female	6	nursery school to grade 4	9999999
214500	male	18	grade 12	0
28863	female	less than 1 year old	n/a or no schooling	9999999
18280	female	11	grade 5, 6, 7, or 8	9999999
40280	male	56	4 years of college	60000
131910	male	9	nursery school to grade 4	9999999
207176	male	28	grade 12	0
50852	female	10	nursery school to grade 4	9999999
42735	male	65	2 years of college	350200

---

## Exercise 3

First, replace all the values of `inctot` that are 9999999 with `np.nan`

```
In [ ]: import numpy as np
sub_acs['inctot'].replace(9999999, np.nan)
sub_acs['inctot']
```

```
Out[ ]: 0          9999999
1           6000
2           6150
3          14000
4          9999999
...
318999      22130
319000      9999999
319001         5000
319002      240000
319003      48000
Name: inctot, Length: 319004, dtype: int32
```

---

## Exercise 4

Calculate the average age of people in our data. What do you get?

```
In [ ]: f'Mean age is ${small_acs["age"].mean()}'
```

```

-----
TypeError                                Traceback (most recent call last)
/var/folders/h9/g02cmrsn6y571zblv96gs56c0000gn/T/ipykernel_10859/1239298680
.py in <module>
----> 1 f'Mean age is ${small_acs["age"].mean()}'

~/miniforge3/lib/python3.9/site-packages/pandas/core/generic.py in mean(self
f, axis, skipna, level, numeric_only, **kwargs)
   10749         )
   10750         def mean(self, axis=None, skipna=None, level=None,
numeric_only=None, **kwargs):
> 10751             return NDFrame.mean(self, axis, skipna, level,
numeric_only, **kwargs)
   10752
   10753         setattr(cls, "mean", mean)

~/miniforge3/lib/python3.9/site-packages/pandas/core/generic.py in mean(self
f, axis, skipna, level, numeric_only, **kwargs)
   10367
   10368     def mean(self, axis=None, skipna=None, level=None, numeric_only
=None, **kwargs):
> 10369         return self._stat_function(
   10370             "mean", nanops.nanmean, axis, skipna, level,
numeric_only, **kwargs
   10371         )

~/miniforge3/lib/python3.9/site-packages/pandas/core/generic.py in _stat_fu
nction(self, name, func, axis, skipna, level, numeric_only, **kwargs)
   10352         name, axis=axis, level=level, skipna=skipna,
numeric_only=numeric_only
   10353     )
> 10354     return self._reduce(
   10355         func, name=name, axis=axis, skipna=skipna, numeric_only
=numeric_only
   10356     )

~/miniforge3/lib/python3.9/site-packages/pandas/core/series.py in _reduce(s
elf, op, name, axis, skipna, numeric_only, filter_type, **kwds)
   4381         if isinstance(delegate, ExtensionArray):
   4382             # dispatch to ExtensionArray interface
-> 4383             return delegate._reduce(name, skipna=skipna, **kwds)
   4384
   4385         else:

~/miniforge3/lib/python3.9/site-packages/pandas/core/arrays/_mixins.py in _
reduce(self, name, skipna, **kwargs)
   258         else:
   259             msg = f"'{type(self).__name__}' does not implement redu
ction '{name}'"
--> 260             raise TypeError(msg)
   261
   262     def _wrap_reduction_result(self, axis: int | None, result):

TypeError: 'Categorical' does not implement reduction 'mean'

```

*Problematic.*

---

## Exercise 5

We want to be able to calculate things using age, so we need it to be a numeric type. Check all the values of age to figure out why it's categorical and not numeric. You should find two problematic categories.

```
In [ ]: age_list = sub_acs['age'].tolist()
        print(age_list)
```

- The two problematic values I find are: "less than 1 year old" and "90 (90+ in 1980 and 1990)"
- Should be strings, i.e. not numeric

---

## Exercise 6

In order to convert age into a numeric variable, we need to replace those problematic entries with values that pandas can later convert into numbers. Pick appropriate substitutions for the existing values and replace the current values. Hint 1: Categorical variables act like strings, so you might want to use string methods! Hint 2: Remember that characters like parentheses, pluses, asterices, etc. are special in Python strings, and you have to escape them if you want them to be interpreted literally!

```
In [ ]: sub_acs["age"] = sub_acs["age"].replace("90 (90+ in 1980 and 1990)" , "90")
        sub_acs["age"] = sub_acs["age"].replace("less than 1 year old", "0").copy()
        sub_acs = sub_acs.copy()
```

```
/var/folders/h9/g02cmrsn6y571zblv96gs56c0000gn/T/ipykernel_10859/1377460654
.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
sub_acs["age"] = sub_acs["age"].replace("90 (90+ in 1980 and 1990)" , "90")
).copy()
```

```
/var/folders/h9/g02cmrsn6y571zblv96gs56c0000gn/T/ipykernel_10859/1377460654
.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
sub_acs["age"] = sub_acs["age"].replace("less than 1 year old", "0").copy()
()
```

## Exercise 7

Now convert age from a categorical to numeric.

```
In [ ]: sub_acs["age"] = pd.to_numeric(sub_acs["age"])
sub_acs.dtypes
```

```
Out[ ]: sex          category
age           int64
educ          category
inctot        int32
dtype: object
```

## Exercise 8

Let's now filter out anyone in our data whose age is less than 18. Note that before made age a numeric variable, we couldn't do this!

```
In [ ]: print("Subsetting the dataset for the people who are above 18 years old:")
adult = sub_acs.loc[sub_acs["age"] > 18]
adult
```

Subsetting the dataset for the people who are above 18 years old:

```
Out[ ]:
```

	sex	age	educ	inctot
2	male	63	4 years of college	6150
3	female	66	grade 12	14000
5	male	50	grade 12	50000
6	male	82	1 year of college	27100
10	male	47	n/a or no schooling	18000
...	...	...	...	...
318998	female	31	5+ years of college	0
318999	female	33	4 years of college	22130
319001	male	20	grade 12	5000
319002	male	47	5+ years of college	240000
319003	male	33	5+ years of college	48000

248538 rows × 4 columns

## Exercise 9

Create an indicator variable for whether each person has at least a college degree called `college_degree`

```
In [ ]: print("Check values again for education variable")
        adult['educ'].value_counts()
```

```
Out[ ]: Check values again for education variable
grade 12                                89743
4 years of college                      47212
1 year of college                      38384
5+ years of college                    29801
2 years of college                     20731
grade 5, 6, 7, or 8                   5942
grade 11                              4763
grade 10                              3942
n/a or no schooling                   3627
grade 9                               3105
nursery school to grade 4             1288
Name: educ, dtype: int64
```

```
In [ ]: print("Create Boolean indicator")
        adult['college_degree'] = (adult['educ'] == '4 years of college') | (adult
        adult['college_degree']
        adult
```

Create Boolean indicator

```
/var/folders/h9/g02cmrsn6y571zblv96gs56c0000gn/T/ipykernel_10859/1651839132
.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
adult['college_degree'] = (adult['educ'] == '4 years of college') | (adult['educ'] == '5+ years of college')
```

```
Out[ ]:
```

	sex	age	educ	inctot	college_degree
2	male	63	4 years of college	6150	True
3	female	66	grade 12	14000	False
5	male	50	grade 12	50000	False
6	male	82	1 year of college	27100	False
10	male	47	n/a or no schooling	18000	False
...	...	...	...	...	...
318998	female	31	5+ years of college	0	True
318999	female	33	4 years of college	22130	True
319001	male	20	grade 12	5000	False
319002	male	47	5+ years of college	240000	True
319003	male	33	5+ years of college	48000	True

248538 rows × 5 columns

## Exercise 10

Let's examine how the educational gender gap. Use `pd.crosstab` to create a cross-tabulation of sex and college\_degree. `pd.crosstab` will give you the number of people who have each combination of sex and college\_degree (so in this case, it will give us a 2x2 table with Male and Female as rows, and college\_degree True and False as columns, or vice versa).

```
In [ ]: pd.crosstab(adult['sex'], adult['college_degree'])
```

```
Out[ ]:
```

	college_degree	False	True
sex			
male	83581	36181	
female	87944	40832	

## Exercise 11

Counts are kind of hard to interpret. `pd.crosstab` can also normalize values to give percentages. Look at the `pd.crosstab` help file to figure out how to normalize the values in the table. Normalize them so that you get the share of men with and without college degree, and the share of women with and without college degrees.

```
In [ ]: pd.crosstab(adult['sex'], adult['college_degree'], normalize='columns')
```

```
Out [ ]: college_degree    False    True
          sex
          male  0.487282  0.469804
          female 0.512718  0.530196
```

## Exercise 12

Now, let's recreate that table for people over 40 and people under 40. Has the difference between men and women in terms of getting a college degree improved, stayed the same, or worsened?

```
In [ ]: below_acs = adult.loc[adult['age'] < 40]
pd.crosstab(below_acs['sex'], below_acs['college_degree'], normalize='columns')
```

```
Out [ ]: college_degree    False    True
          sex
          male  0.532658  0.437058
          female 0.467342  0.562942
```

Comparing the overall sample to the subset of those below 40, on the first look the gap has overall stayed pretty much the same. However, marginally, the majority attribution for those without a degree have changed:

- We find that in the full sample, 51% (i.e. the majority) of those who did not have a degree were women, while in the subset the majority without a degree are men (53%). When looking at those with a degree the majority attribution has stayed the same but distribution has changed:
- In the full model, 53% of those with a degree were female, in the subset that number increases to 56%.



```
In [ ]: above_acs = adult.loc[adult['age'] > 40]
pd.crosstab(above_acs['sex'], above_acs['college_degree'], normalize='column')
```

```
Out[ ]: college_degree    False    True
sex
male    0.464684  0.487035
female  0.535316  0.512965
```

Comparing the overall sample to the subset of those above 40, on the first look the gap has overall stayed pretty much the same. THE majority attribution has stayed the same but distribution has changed:

- We find that in the full sample, 51% (i.e. the majority) of those who did not have a degree were women, while in the subset that number increases to 53%.
- In the full model, 53% of those with a degree were female, in the subset that number decreases to 51%.

## Conclusion

As we see the age gap is generally not very strong but it is interesting to see how the age gap behaves across generations. We find differences in the majority attribution for those above and below 40. On a first glance we can state that women are more likely than men to have a degree in the age group below 40. This would require further investigation.